


# OS bypass performance without OS bypass

Usenix 2010

# Outline

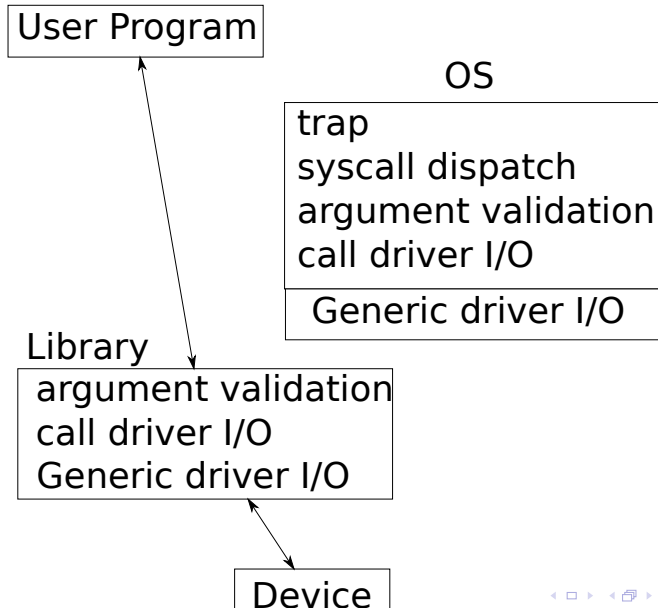
## 1 Currying and process-private system calls

# Standard path

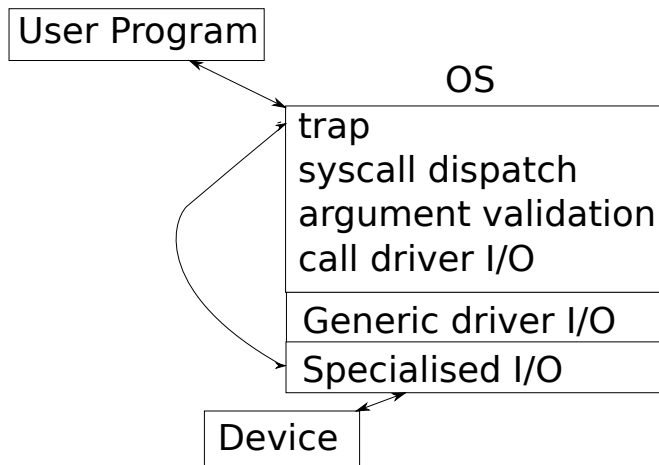


- OS overhead known to be primary limit even with very efficient devies

# OS Bypass



# Pushing IO back into the kernel via currying and process-private system calls



# How it is used

- `fd = open( "/dev/torusctl")`
- `newsyscall = write(fd, "fastpath pointer size");`
- parameters and operation "curried" into new system call
- `newsyscall` attached to this process and its children only
- `newsyscall()`; – no parameters, fast path
- no parameters and bypasses all system call checking
- 700 ns measured time on barrier
- Now any program, not just mpi programs, can use the network
- Being integrated into new mpi library now