

XCPU3: Workload distribution and aggregation

Pravin Shinde

Parallel and Distributed Systems(PDCS)
Vrije Universiteit, Amsterdam
Email: pse220@few.vu.nl

Eric Van hensbergen

IBM Research Lab
Austin, USA
Email: ericvanhensbergen@us.ibm.com

Abstract—As clusters and grids are reaching to domains like commercial application, it is becoming more clear that the existing workload distribution solutions which are tailored for high performance scientific applications, are not efficient for dynamic data-flow applications and many-task computing (MTC).

Most of the existing workload distribution solutions that concentrate more on HPC applications and are not efficient in quickly deploying the dynamic workload involving large number of small jobs which is the typical case in dataflow applications. Also developing and deploying applications on existing infrastructures still is a non-trivial task requiring special runtimes, middleware support and/or language dependence.

We have explored an alternate approach for simplifying workload distribution and aggregation based on synthetic filesystems and the private namespace concepts of Plan 9. Instead of sending workloads to a remote cluster, XCPU3 works by bringing the remote cluster to the user's desktop, leading to full control over the compute environment. XCPU3 provides a filesystem interface, making it runtime and language agnostic. It also allows executing multiple jobs simultaneously in isolation, leading to better resource utilization. XCPU3 makes all nodes independent and equivalent in functionality leading to the ability to localize the decision making and ability to handle dynamic workloads. XCPU3 provides easy to use, flexible and scalable infrastructure for workload distribution and aggregation that can be used by dynamic workloads and dataflow applications.

I. INTRODUCTION

Due to the limits on the silicon technology, the trend is moving in the direction of multicore and manycore architectures. Having a large number of less powerful computers is assumed to be economically more viable than having a single more powerful computer. The increasing acceptance of clusters and grids coupled with the emergence of cloud computing and availability of compute resources as utility via [1] [5] [4] has increased the need for parallel programming and workflow management.

The initial development around parallel programming and the traditional infrastructure around clusters and grids is designed to suit the requirements of the scientific community which tend to run large compute intensive applications for long durations. With growing use of clusters and grid, there is increasing interest in deploying new genre of applications like dataflow applications and many-task applications. Typical commercial applications fall into the category of dataflow applications and involve executing lot of small and interdependent applications for small amount of time. Unfortunately, most of the existing workload deployment solutions are ill-

suited for these applications. These solutions are efficient in deploying large compute-intensive application which will run for long duration, but are highly inefficient in deploying large number of small applications for small amount of time.

The major obstacle in widespread adoption is the expertise needed for writing parallel applications. Also, dependence on a specific language or runtime environment for parallelism dictates the need for major re-writing all existing applications for exploiting the parallelism.

Another crucial property exhibited by dataflow applications is that the amount of resources needed for total computation is unpredictable in the beginning as it may depend on results of intermediate computations. This leads to the requirement of dynamically allocating new resources whenever needed. Most of the current job schedulers do not support such adaptability and flexibility.

In this paper, we target the problem of creating a scalable infrastructure with the flexibility needed for running dataflow applications. We primarily target the Blue Gene [26] super-computer as our testbed as it allows us to test our infrastructure on up to 64K nodes.

The challenge is to provide the flexibility needed by dataflow applications using a very simple interface. We need to provide an infrastructure with support for dynamically adjusting the resource reservation which can help in coping with unpredictable resource requirements. This infrastructure needs to provide quick job-startup time even for large numbers of jobs, allowing overall good performance for dataflow applications. It also needs to provide a simple interface for communication between nodes even when those nodes cannot communicate directly with each other. We do not want the application development for this infrastructure to be constrained to a particular language, runtime, middle-ware or platform. And we need to run most of the existing applications without modification.

Our primary aim is not performance improvement, but to present a new and elegant way for workload distribution and aggregation. It is aimed at demonstrating that simple filesystem interfaces and a simple design based on the basic principles of the Plan 9 [23] distributed operating system can be easy, intuitive and scalable to a large number of nodes.

This paper is organized as follow. The next section presents the background work which has inspired this project. The *Design section* will discuss the design and the reasoning behind it. *Implementation section* will discuss how system is

implemented. The *Filesystem interface* section explains the use of synthetic filesystem and presents few examples to show how easily this interface can be used. Evaluations section presents some results showing that this approach is feasible. Related work section positions XCPU3 in the context of other related projects. Conclusion section will conclude this project.

II. BACKGROUND

This section discusses the the previous work from which XCPU3 is evolved.

A. Plan 9

Plan 9 is an operating system which was specifically designed for dealing with networked systems. It is based on principles like *everything is file* and *private namespaces*[24]. It provide applications ability to build custom namespace by mounting and binding remote resources and then run in network oblivious manner. The usefulness of these features can be seen in CPU[2] which is Plan 9's way for using remote compute resources. Unlike Unix based remote login and SSH where remote processes run in the namespace and environment of remote node, CPU runs the remote processes in local namespace and environment. One of the key advantages of this model simplifies the remote workload deployment as even though workload runs on remote resources, it runs in local environment without getting affected by the environment on remote node.

B. XCPU

XCPU[21] [20] was an attempt to bring the functionality of CPU into mainstream operating systems like Linux. XCPU allows process management using a **filesystem interface**. As this underlying psudo filesystem can be exported over network to other nodes, it simplifies the sharing of the compute resources. XCPU does not provide private namespace, which is needed for transparantly accessing remote filesystem. XCPU works around this limitation by pushing the executables and it's dependencies needed for requested remote execution.

To achieve scalable pushing of programs and their dependencies, XCPU uses a few compute nodes for distributing the programs and dependencies during the program startup time creating a tree hierarchy of the nodes. This is known as the **tree spawn mechanism** and it improves the scalability for XCPU to larger numbers of nodes. With automatic pushing and the tree spawn mechanism, XCPU not only simplifies the deployment of applications but it also scales them better with quicker job startup time.

C. XCPU2

XCPU2 [19] is evolved from XCPU and is aimed to give more flexibility by using *private namespace* support added recently into Linux kernel[13]. This allows remote processes to execute in the namespace which provides local filesystem view. By controlling the environment and filesystem view of all the remote processes, every developer is free to use any libraries and development setup, regardless of what is installed

on remote nodes. As these processes run in the isolation provided by the private namespace, they do not pollute the environment and filesystem view of other processes. This allows executing multiple processes of different users on same remote node without conflicting with each other.

XCPU2 makes clear distinction between various nodes which are assigned with specific responsibility. The nodes are classified into *control node* which is responsible for reserving the compute nodes, *job control node* which are responsible remote job management and namespace management on compute nodes. The *compute nodes* are responsible for actual computation by running the processes on inside custom namespace on users behalf.

The XCPU2 lacks fine-grained granularity of the control. It only supports managing all *compute nodes* together in identical way, and does not give any mechanism to control each *compute node* individually. This lack of control on every node is a major deterrent for dataflow applications. This limitation was major motivation for us to develop XCPU3 for providing much needed flexibility and fine grained control over over all compute nodes using simple interface.

III. DESIGN

The key challenge in this design was to provide the flexibility needed by dataflow applications while achieving scalability to the large number of nodes available.

A. Key design decisions

We decided to adopt a few guiding principles which will help us to reach our goal of flexibility with scalability. We also decided to keep the design simple, and we believe that the simple design should lead to the simple and flexible interface. We also hope that the lack of complexity may help in improving scalability. In this section we present those decisions which influenced most aspects of system design and implementation.

1) *Localization of decision making*: The key requirement for us was scalability to a large number of nodes. We planned to design the system without any central component which should have knowledge of the entire system. We plan to distribute and localize the computations as well as decisions like scheduling, job management and workload distribution/aggregation.

We avoid decision making based on global knowledge and promote use of local information. We hope that if each node tries to attempt a local optimum, we will reach the global optimum. This may not be true in all cases, We hope that in those cases, we hope to perform acceptably well if not optimal.

2) *Independence of every node*: As we plan to distribute and localize all functionalities, it was essential to replicate these functionalities at multiple levels making localization possible. The granularity of functionality replication decides the granularity of control, and hence influences the flexibility provided by the system. As we aim for maximum flexibility, we have decided for replicating the following three functionalities at each node.

- 1) **Resource reservation:** Each node should be able to reserve more resources on its own without involvement of any central entity.
- 2) **Job management:** Every node should be capable of starting and managing new jobs using his reservations.
- 3) **Computation:** Every node should be able to perform the computation by running the requested application in isolation and returning the results.

We intend to provide each node with the capability to perform all of these roles simultaneously instead of binding them in one role at one time. This design makes the interface to every node a building block identical to each other, and provides the flexibility to build any structure with these building blocks.

There are a few downsides in making every node independent. With independence of every node, each node can be a source of failures and faults. One will need to come up with better ways to deal with faults and failures when so many sources of them are present. This takes the XCPU3 in realms of **Distributed Systems** opening many more possibilities and questions.

For purpose of this exploration, we avoid these complexities by assuming a very simple model for handling failures. Any failure anywhere in the system will result in the failure of the entire operation. We assume that failures will be infrequent, so aborting and restarting operation should be acceptable for such infrequent failures.

3) *Filesystem Interface:* We want to keep XCPU3 interface agnostic from language, runtime and middleware. Plan 9 has proven that the filesystem interface is very flexible and yet powerful in the world of distributed applications. We aim to follow the same principle of **Everything is a file** from Plan 9.

Every node will provide access to its services via filesystem interfaces. This interface will be exported as the filesystem over the network so that other nodes can use it. Other nodes can mount this filesystem and use it as interface for interacting with that node. Multiple remote nodes can be aggregated into the filesystem hierarchy providing a clean and easy way to access them.

Multiple overlay views can be created by *binding* the same filesystem at multiple locations with different names. This ability of creating ad-hoc overlays allow users to arrange remote resources as per his needs without worrying about their actual locations.

Other advantages of using filesystem interfaces are that

- 1) Existing tools/commands used with traditional filesystem can be directly used with XCPU3.
- 2) filesystems come with their own mechanism for access control list for providing the security. We can leave the security to these already proven mechanisms instead of implementing our own.
- 3) We inherit the ability to export, mount and bind the filesystem without writing any explicit code for it.
- 4) Filesystem interfaces are simpler to program than socket interfaces. This can lead to simpler code and hence lesser

bugs.

- 5) Users don't need special privileges or administrative access to interact with the filesystem. This simplifies the user experience in running XCPU3 based applications.

The filesystem interface has following limitations which affect our system design.

- 1) The critical limitation concerns the POSIX standard for failure reporting in filesystem. POSIX standard dictates that file operations should report their success or failure in the form of a single number which may not be enough to provide the information about the reason behind failures. Plan 9 overcomes this limitation by returning a string which can provide more verbose information instead of a single number. But this breaks the POSIX compatibility. So, we need to find an alternate way to report information about failures in POSIX compatible way.
- 2) Another drawback is the way a failure of remote services is detected. The filesystem interface relies on the underlying networking protocol for detecting failures by waiting for timeouts, and then reporting them back to users as an error. This makes the filesystem interface less desirable where quick failure detection and recovery is needed.

Our decision to choose the filesystem interface despite of its drawbacks is the trade-off we are willing to do for flexibility and simplicity. We limit ourselves with the assumption that failures will be infrequent. With this assumption, we are willing to accept the delays in reporting failures at the remote end.

IV. IMPLEMENTATION

This section discusses the implementation issues in more detail. This chapter concentrates on how XCPU3 was implemented and discusses implementation decisions which have proven to be important.

A. *Inferno*

Inferno[28] is an open source distributed operating system which is a direct descendant of the **Plan 9** operating system. It runs natively on multiple hardware platforms and can also run as a user-space application on top of other operating systems. For the purpose of this document "*hosted Inferno OS*" or "*userspace Inferno*" refers to Inferno running as user-space application on top of other operating systems.

As we are aiming for a flexible heterogeneous environment with different operating systems and different architectures, we found the Inferno operating system to be an attractive platform. It allowed us to develop XCPU3 once and easily deploy it on different platforms while enjoying the Plan 9 features on all those platforms.

This decision does come with the cost of performance loss. Each node needs to run hosted Inferno in user-space which takes up some resources of the node. Our choice of Inferno gave us the flexibility and features of the Plan 9 operating system on other operating systems. It allowed us to quickly

implement and deploy our filesystem on multiple platforms easily at the cost of some performance.

B. Implementation details

XCPU3 is implemented as the filesystem in the Inferno kernel. Figure 1 shows the placement of XCPU3 in the context of applications, host operating system and the Inferno.




Fig. 1. XCPU3 Structure

The XCPU3 filesystem is exported by Inferno using the 9P protocol. This protocol is used by Plan 9 and Inferno extensively to access any file. Recently the Linux kernel has added support for the 9P protocol[18]. This allows Linux to mount filesystems exported over the 9P protocol. Other Unix based operating systems can use FUSE[3] for accessing the 9P based filesystems.

In a typical setup, Inferno runs in the hosted environment on an operating system like Linux or Plan 9. The hosted inferno will export the XCPU3 filesystem which can be mounted by the local host operating system. The applications interact with this mounted XCPU3 filesystem using the native filesystem interfaces (e.g. VFS for Linux). Any interaction with this XCPU3 filesystem is communicated to Inferno using 9P. The XCPU3 kernel module inside the inferno kernel then receives and interprets the user actions. If needed, it uses services from the host operating system or from other XCPU3 filesystems deployed on remote locations. The XCPU3 filesystem then sends the prepared response over 9P. The host operating system will relay this response back to the application via the local filesystem interface.

C. The big picture (connecting multiple XCPU3 nodes)

The real power of the XCPU3 is the ability to connect with each other and form bigger instance of XCPU3. XCPU3 filesystem is exported over 9P providing same functionality and interface to both local and remote applications. This allows XCPU3 instances to connect with other XCPU3 instances and distribute some of their responsibilities among them. The 9P protocol shields the XCPU3 from complications of remote filesystem access and works as transparent glue for connecting multiple XCPU3 instances together.

1) *Central Services*: The ability to configure many XCPU3 nodes into hierarchy is provided by the *central services*. In contrary to the name, central services is highly distributed and every XCPU3 runs an independent instance of the central services.

The central services synthetic file server which provides a simple hierarchy of directory mount points representing remote nodes. Mounts of the remote nodes or binds of previously mounted remote nodes are accomplished within this file system such that anyone who mounts our name space can also see (and access) anyone we have mounted transitively, in such a way a child node can access a parent nodes, other children, or the parents nodes parent and so forth. Any node could establish themselves within a hierarchy by binding a parent's central service directory to the name `[/csrv/parent]` and then tell the parent to back-mount their name space (allowing two way traversal). In this way children register with parents triggering the cross-mounts and establishing a two way link between them. Each XCPU3 instance need to know only the information about its parent and children in the hierarchy and the all XCPU3 nodes initiate these connections leading to the distributed creation of the XCPU3 node hierarchy.

Figure 2 tries to give a simple overview of how this synthetic filesystem view is populated based on the underlying mount connections between the nodes.

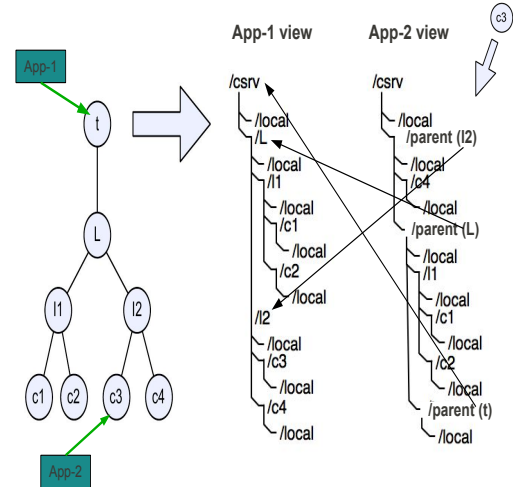


Fig. 2. Sample filesystem interface for sample the topology in XCPU3

Assuming the links between the nodes are created by the remote node mounts in central services, this diagram shows how the filesystem views at different nodes encompasses the whole network, even though each node is only connected to its neighbours. The XCPU3 filesystem starts with the `[/csrv]` directory. The location `[/csrv/local/]` presents the local resources whereas `[/csrv/parent/]` presents the XCPU3 filesystem of the parent node. All other directories in the represents the XCPU3 filesystem of the children nodes. It can be easily seen that both *App-1* and *App-2* have access to all

the nodes even though they are running on different nodes. In this design, every node has to worry about only its children and the parent, other topology falls in the place automatically.

Even though the filesystem view at different nodes differ from each other, all nodes have access to the full topology. The `[/csrv/]` filesystem view encodes enough information within itself that any nodes can construct the global view and figure out their own position within the global view.

Just because every node can construct the global view, does not mean that it must use this global view for making any decision or performing a typical operation. The nodes mostly use only the local view for decision making and operations. This local view includes the parent node and the children nodes.

V. FILESYSTEM INTERFACE

This section discusses the filesystem interface and how it can be used for managing the local and remote resources. Figure 3 gives the the high-level view of the hierarchy of the XCPU3 filesystem. This section will discuss the overall hierarchy and few important files in the XCPU3.

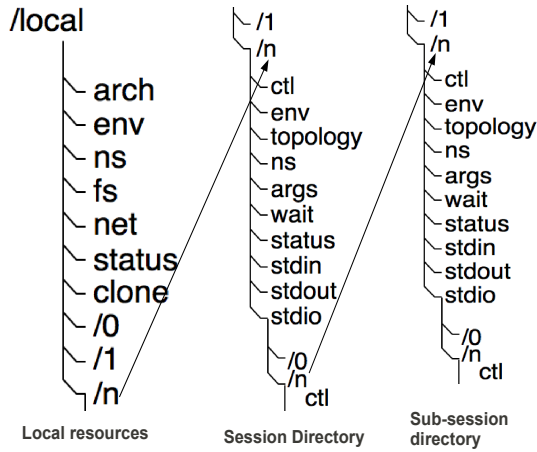


Fig. 3. Filesystem interface in XCPU3

The location `[/csrv/local/]` points to the local resources available on the node and the files present in this directory allows managing these local resource. Few files here provide information about the system, like `[arch]` file tells the architecture and operating system of the node while `[status]` file provides information about total amount of resources available, including the remote resources which are direct or indirect decedents of this node. Files like `[env]` and `[ns]` allowed controlling the default environment and the namespace of the processes created on this node. `[fs]` and `[net]` are links to the local filesystem and networking resources available on the node. The `[clone]` file provides the interface to create new sessions which is the unit of the execution.

The sessions are represented as directories with session-id number as a name. These sessions are self-contained and provide interfaces to manage the execution of the session. Files

like `[env]` and `[ns]` are present in session directory also, and can be used to overwrite the default environment and namespace specifically for this session. File `[ctl]` is used for controlling the execution and file `[stdio]` is used to manipulate standard input and output. Sessions can play one of the two roles between *single execution unit* and *aggregation* of other sessions. Following is the attempt to define these two modes.

```

SESSION -> AGGREGATION | execution
AGGREGATION -> SESSION+

```

The **execution mode** executes the requested program and provides the output. The **aggregation** mode aggregates and manages the multiple sessions which are created as sub-sessions for this session. In other words, aggregation session creates sub-sessions, divide the work among them and aggregate back the output. These sub-sessions are nothing but sessions created by using `[clone]` file on local or remote systems. These newly created sessions are binded as sub-directories inside the current session directory.

A. Workload distribution and aggregation

The workload distribution and aggregation works by using the dual nature of the sessions. The first step in this is reservation, which is responsible for creating the tree of the sessions. Whenever any session receives the reservation request, it evaluates it by considering, if it can execute locally or it should break it and divide between its available children. Depending on the decision, it can play one of the following two roles.

- 1) **Execution session:** When reservation is request is small enough or there are no children nodes available, then session behave as execution session by performing the requested operations locally. The execution sessions are the leaf node in the session tree.
- 2) **Aggregation session:** When request is bigger than what session can handle, it divides the request between children nodes. The session does this by creating new session and sending the divided reservation request to the each child node available in it's **CSRV** hierarchy. These newly created sessions will be binded as directories in the directory of the current session hence making them the sub-sessions of the current session. The current session now acts as in aggregation point for these sub-sessions by passing the any execution request or input to all sub-sessions and merging the output from all sub-sessions to produce the final output. These aggregation sessions constitute the non-leaf nodes in the session tree.

This way, recursive tree of sessions is created whenever new reservation request is received. This tree of the session is then used to request the execution and collect back the results. This design not only distributes the actual execution, but also it distributes the distribution and aggregation process as this work is divided between all aggregation sessions.

B. Examples

It is hard to measure the *ease of use* and *simplicity* factor. So, we are presenting following examples to show how easy it is to use XCPU3 infrastructure. We are giving these examples of using the bash shell in Linux.

1) *Traditional application deployment*: This example presents how the default aggregation behaviour of the XCPU3 can be used to deploy large number of applications.

We assume that the XCPU3 filesystem is mounted on `./mpoint/` directory. This can be done by using FUSE or 9VFS[17].

```
$ less ./mpoint/csrv/local/clone
0
```

The above command is an example of creating a new session. The contents read from the `[clone]` file represent the session-ID. Now we use session 0 for performing actual execution.

```
$ echo "res 4" > ./mpoint/csrv/local/0/ctl
$ echo "exec date" > ./mpoint/csrv/local/0/ctl
$ cat > ./mpoint/csrv/local/0/stdio
Fri May 7 13:53:58 CDT 2010
Fri May 7 13:53:58 CDT 2010
Fri May 7 13:53:58 CDT 2010
Fri May 7 13:53:58 CDT 2010
$
```

The first echo command sends the request for reserving 4 remote resources. The next echo command submits the request for executing the `date` command. And the `cat` command on `[stdio]` returns the aggregated output to the user. This example shows all the complexities about finding, connecting and using the remote resources is hidden behind the filesystem interface. This approach can be used in the *trivially parallelizable applications* where the same application is deployed on all the nodes.

2) *Dataflow deployment*: In this mode the role of XCPU3 is limited to the reservation. An application is given the capability of doing the workload distribution and aggregation with the help of the filesystem interfaces provided by the XCPU3. Instead of interacting with the aggregation points, these dataflow applications can interact directly with the sessions responsible for the actual execution using the filesystem hierarchy of that session.

In this example, we will try to create a small pipeline of two commands `date | wc`. But we will create this pipeline across multiple nodes. The objective of this example is to show how the underlying complexities are hidden from the application.

Lets assume that session 0 is created by opening `[clone]` file as shown in the previous example. The following commands will create the desired pipeline.

```
$ echo "res 2" > ./mpoint/csrv/local/0/ctl
$ echo "exec date" > ./mpoint/csrv/local/0/01/ctl
$ echo "exec wc" > ./mpoint/csrv/local/0/01/ctl
```

```
$ echo "xsplice 0 1" > ./mpoint/csrv/local/0/ctl
$ cat ./mpoint/csrv/local/0/1/stdio
1 6 29
$
```

The first command `[exec date]` is sent to 0'th sub-session and the second command `[exec wc]` is sent to the 1st sub-session. The `[xsplice 0 1]` request tells the parent session to redirect the output of the 0'th session to the input of the 1st session. The `xsplice` command can be seen as a pipe operator of the shell script for redirecting the output of one command to the input of other command.

The above example is equivalent of executing `date | wc` on the shell, but with the difference that both commands are executed on a different remote machines while sharing the same namespace.

The XCPU3 infrastructure relies on userspace applications like the **PUSH shell**[16] for creating more complicated DAG and mapping them on underlying XCPU3 infrastructure.

VI. EVALUATION

This section presents the evaluation of the XCPU3 infrastructure from the perspective of quicker deployment of a large number of small jobs while giving clean interfaces and abstractions.

A. Experimental setup

We have performed our evaluations on a Blue Gene setup with 512 nodes. This setup is visually presented in a figure 4. We run hosted Inferno on all the compute nodes, IO nodes and the controller node. The user interacts with the XCPU3 instance on the controller node for job submission.




Fig. 4. Setup for evaluation

1) *Scalability of XCPU3*: Our first objective is to show how quickly we can do deployment and execution of a large number of small applications. We have avoided using larger applications as the bigger runtimes of larger applications tend to amortize the overhead in the deployment of the application. We have used the `date` command as the application for deployment. This is a small application and does not need any external inputs and produces small output. Each deployment involves session creation, reservation, execution, output aggregation and termination of the session. We deployed varying

numbers of execution of this application on the cluster of 512 nodes. The number of requested executions increased exponentially from 1 to 2048 executions.

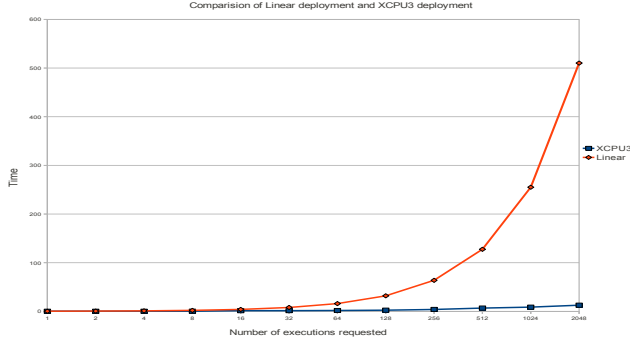


Fig. 5. Comparison if XCPU3 with sequential deployment

Figure 5 gives an initial perspective of how XCPU3 performs relative to sequential performance. This graph plots the total time taken by XCPU3 and the hypothetical time it may take for performing the same amount of work on one machine. This graph shows that the XCPU3 is successfully able to exploit the parallelism for deploying the jobs quickly. The XCPU3 deploys 2048 jobs in 12.66 seconds whereas sequential execution would take upto 510 seconds. In the graph, the line showing sequential scaling looks exponential, but that is because number of requested executions increase exponentially.

2) *Job deployment without input*: This section aims to further analyze the performance of XCPU3. Again we are concentrating on similar deployment. We have recorded the time taken by each of the following stages in the deployment on the XCPU3 infrastructure.

- 1) Reservation: Create a new session, and request the reservation by writing `res n` into the session `[ctl]` file. Here `n` varies from 1 to 2048 representing the number of executions requested.
- 2) Execution: Request the execution by writing `exec date` into the session `[ctl]` file.
- 3) Aggregation: Collect the output generated by all the executions by reading the session `[stdio]` file.
- 4) Termination: Closing all the files and terminating the session.
- 5) Housekeeping: Additional time taken before, between and after above steps.

Every deployment starts with the creation of the session followed by the reservation, execution, aggregation and then ending with termination of the session. We have taken the average value over multiple runs for our analysis.

Figure 6 presents the results of deployment of the date command in the form of graph. This graph presents the breakup time for various stages of the deployment using the XCPU3 infrastructure.

From this graph we can observe that the session termination and the housekeeping overheads are negligible compared to

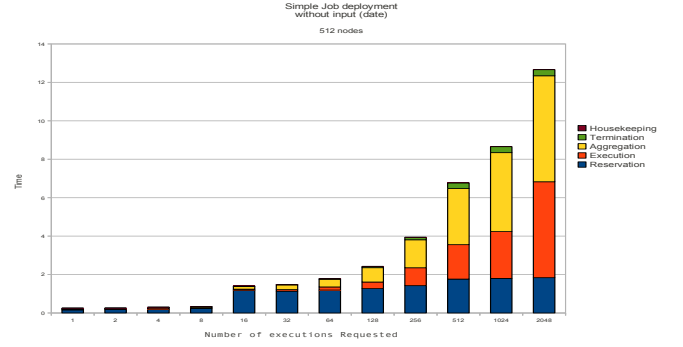


Fig. 6. Deployment without input

the time taken by reservation, execution and aggregation. So, we can ignore these two overheads in our future evaluations. For jobs of up to 128 deployments, the reservation time dominates everything else. But for larger numbers of deployment execution and aggregation time increases rapidly while reservation time remains relatively constant. This shows that reservation time is not directly dependent on the number of deployments, whereas execution and aggregation time are directly proportional to the number of deployments.

Now, let us try to analyze why reservation time is independent of the number of deployments. The reservation process involves traversing the underlying topology tree of nodes till the reservation requirements are satisfied. All the children on the same level are traversed parallelly at the same time. This way, each level is traversed in the constant time, independent of number of nodes in that level. Another aspect of the reservation mechanism which helps here is that the amount of data written and read from the `[ctl]` file and the amount of data exchanged between nodes for communicating reservation request is fixed in size and independent of the number of deployments requested. With these two properties, the reservation time becomes directly proportional to the depth of the tree and not with the number of nodes.

We can observe the above relation in figure 6. The reservation time remains relatively constant for deployment requests from 1 to 8. Then it sharply increases between 8 to 16 and remains almost constant for all the requests between 16 to 2048. This can be attributed to underlying cluster topology. Figure 4 shows the presence of the 8 IO nodes in the first level. This enables satisfying the requests which are smaller than 8 executions. For larger requests, one more level needs to be traversed in the topology, introducing delays. The time taken for reservation remains almost constant between 16 and 2048 executions as all these reservation requests essentially traverses the same depth. We can conclude from these observations that *the time taken for the reservation is directly proportional to the depth of the tree*.

Now let us discuss, why the same property is not exhibited by execution time or aggregation time. We have discussed in the implementation chapter that all read and write requests are performed in parallel between all the nodes in the same

level. But the amount of data exchanged for aggregation and execution is not constant. This data is directly proportional to the number of nodes involved. With the increase in the number of requested deployments, the amount of data to be exchanged also increases, leading to larger aggregation time. The execution time is also similarly affected as all compute nodes will try to fetch the binary of the executable from the initiating node leading to the copy of the data. These observations lead us to the conclusion that *the time taken for the execution and aggregation is directly proportional to the number of deployments requested*.

3) *Job deployment with input*: Our next evaluation involve the deployment of an executable `wc` which needs input. This command counts the number of lines, words and characters in the input file. This is an interesting case for our infrastructure as this deployment involves the distribution of inputs to all the sessions. This introduces a new stage in the deployment process in addition to the 5 stages we described in the above section. This stage will be the **input** stage and involves distributing the input data to all the sessions which are responsible for execution. By default XCPU3 will broadcast the input to all the compute nodes, but we have plans to introduce filters in the PUSH shell which can partition the input given to the compute nodes.

Figure 7 presents the results of evaluations involving the distribution of the input.




Fig. 7. Deployment with input

These results enforce our observations that reservation time is directly proportional to the the depth of the tree whereas aggregation and execution time are directly proportional to the number of deployments requested. In addition to these observations, we can also observe that the input time exhibits behavior similar to the execution and aggregation time. This observation can be attributed to the fact that input distribution implementation is similar to the output aggregation implementation. And also the amount of data to be exchanged for input distribution depend on the number of deployments requested as this data needs to reach all sessions responsible for execution. This increases the amount of data to be exchanged with any increase in the number of the deployments requested. We can conclude with this observation that *the input time is directly proportional to the number of deployments requested*

B. Dataflow workloads

The evaluations presented in this chapter helps us in understanding how fast XCPU3 can deploy small jobs and aggregate the results produced by them. Now let us relate how these observations can justify our claims about dataflow applications. Typical deployments of the dataflow applications are similar to the above experiments as it involves starting up a large number of small jobs. But the similarity is over at this point. Dataflow deployment does not always involve the input distribution and output aggregation stages. These deployments work by feeding the output of one computation as input for other computation. At the end of the computation, a user needs to read the output of only selected compute sessions which does not need any aggregation. What we can conclude from the above description is that XCPU3 performs really well in the stages like reservation which are important for dataflow deployments. The stages like input distributions are output aggregation where XCPU3 is relatively slow are not needed by dataflow deployments. This makes XCPU3 ideal for rapid deployment of dataflow workloads.

Unfortunately we do not have concrete measurements and evaluations to back our predictions. XCPU3 is one of the piece in the envisioned solution for problem of efficient and easy deployment of large-scale dataflow applications. We are still working on the integration of userspace applications like PUSH[16] with XCPU3 which will further simplify the dataflow deployment.

XCPU3 is an infrastructure which provides the needed flexibility, speed and ease of use for dataflow workloads. This chapter demonstrates the speed that can be achieved. It is difficult to measure the properties like *ease of use* and *flexibility* but the examples presented in the filesystem interface chapter should give some insights of all the possibilities opened up by the XCPU3.

VII. RELATED WORK

This section tries to put XCPU3 in the context of other work which has been done in this area. We will briefly discuss the various related projects and how they differ from XCPU3.

A. Historical solutions

The concept of workload distribution between multiple machines was introduced in the research by research projects like Amoeba [25] and Cambridge distributed computing system[22]. These were developed as a general purpose distributed operating system which can solve the problem of scarcity of compute resources. These systems introduced many useful concepts, but easy availability of PC made these systems outdated.

B. Traditional solutions

As we are approaching the limits on silicon technologies the trend has changed to design of multicore and manycore chips and increased the use of clusters and grids. These clusters and grids had independent operating system on each machine and were loosely coupled using networking. This lead to the

remote process management based on remote login and remote shells like SSH[11]. Most of the existing middleware and workload distribution systems use the above approach with wrappers to simplify the user interface.

The limitation of using SSH is that the application has to run in the namespace of the compute node instead of client node. And as SSH connections are one-to-one, they do not scale well for managing large number of remote nodes. Attempts like Parallel SSH (PSSH)[10] to overcome this limitation of scalability but at cost of losing the flexibility of fine grained control.

C. BProc

BProc: Beowulf Distributed Process Space [14] is developed as an alternative to the loosely coupled SSH based job deployment. BProc works by tightly coupling the kernels running on cluster nodes and provides distributed process ID space which allows a front node to locally control all remotely started processes.

Bproc is definitely a step towards more control on remote processes but needs modified kernel on all systems and remote processes still run in remote namespace instead of the client namespace. These restrictions lead to the losses of the flexibility which is needed for dynamic workloads.

D. Multicast Reduction Network (MRNet)

The MRNet[6] is designed for efficient multicast and data aggregation. This project concentrates on optimizing group communication but it is language dependent. The MRNet API is in C++ and one needs to develop the applications using this API leading to lot more efforts from the developer side to use this infrastructure in comparison to XCPU3.

E. Streamline

The idea of using filesystem interface for data-streams is also used by the *pipefs*[15] which is part of the streamline[12] project. This approach of using filesystem interface is quite similar to XCPU3, but the objectives of both these projects are different. Pipefs concentrates on faster I/O with filesystem interface for the Linux kernel while XCPU3 concentrate on the flexible workload management.

F. Dryad

Dryad[29] is a distributed engine for data-parallel applications which is designed with a primary focus on simplicity of the programming model, reliability, efficiency and scalability. Dryad is well suited for dataflow applications where amount of computation needed is already known. But it does not provide way to dynamically adjust the reservations to deal with changing workloads.

Dryad addresses the issues like fault tolerance which are currently ignored by XCPU3, but XCPU3 is more flexible as it can support dynamic workload and changing DAG of dataflow deployments. Also the filesystem interface of the XCPU3 is much simpler and cleaner compared to the class based C++ interface of Dryad.

G. Condor

Condor[27] is *high-throughput distributed batch computing* system which exploits *opportunistic computing* and idle CPU cycles of voluntary workstations for better performance. Condor works by breaking the large problem into smaller tasks and submitting these smaller tasks to voluntary workstations. Condor handles scientific workflow application by using the meta-scheduler DAGMan (Directed Acyclic Graph Manager) which allows user to specify the various dependencies within tasks in form of directed acyclic graph.

Unfortunately, the voluntary model for condor restricts the direct communication between the workstations for security reasons. They typically communicate only with the agent responsible for submitting the job. This lack of the ability to communicate directly with other computational nodes limit the usefulness of Condor for communication intensive dataflow applications.

H. Other commercial solutions

There are few commercial Workload deployment solutions available. The *Oracle Grid Engine*[7] is a batch-queuing system which can accept, schedule, dispatch and manage the remote executions on clusters. Other solutions involve *PBS Works*[8] which is workload and resource management solution and *platform LSF (Load Sharing Facility)*[9] batch job scheduler which is aimed to help in scheduling jobs on private clouds. Most of these solutions target to simplify the resource management and job submission/scheduling, but they do not try to simplify the communication issues within applications.

VIII. FUTURE WORK

The XCPU3 is an infrastructure which we plan to extend it further to support fault tolerance and add capability to selectively restarting the sub-sessions for environment where faults are more frequent.

As XCPU3 is filesystem interface and can be implemented by any filesystem, we plan to implement this interface natively on traditional operating systems to improve the performance by avoiding the virtualization layer added by the Inferno.

IX. CONCLUSION

XCPU3 project was initiated with the goal of creating alternate job deployment mechanism primarily for manytask and dataflow deployments which can deal with dynamic workloads. We have kept the design simple with the assumption that simple design leads to the better scalability and flexibility. We have tried to show that from examples that the XCPU3 interfaces are easy to use in language and runtime independent fashion. And with evaluation section we have tried to prove that the XCPU3 is viable option for scalable deployment of many-task and dataflow workloads.

REFERENCES

- [1] Amazon elastic compute cloud.
- [2] Cpu man page. *Plan 9 Programmer's Manual*.
- [3] Fuse: Filesystem in userspace.
- [4] Google app engine.

- [5] Microsoft azure services platform.
- [6] Multicast reduction network (mrnet).
- [7] Oracle grid engine.
- [8] Pbs works.
- [9] Platform lsf(load sharing facility).
- [10] Pssh: Parallel ssh.
- [11] Rfc 4251: Ssh.
- [12] The streamline.
- [13] Eric Biederman. Multiple instances of the global linux namespaces. pages 101–112, 2006.
- [14] Sung-Eun Choi, E.A. Hendriks, R.G. Minnich, M.J. Sottile, and A.J. Marks. Life with ed: a case study of a linux bios/bproc cluster. pages 33–39, 2002.
- [15] Willem de Bruijn and Herbert Bos. Pipesfs: fast linux i/o in the unix tradition. *SIGOPS Oper. Syst. Rev.*, 42(5):55–63, 2008.
- [16] Noah Evans and Eric Van Hensbergen. Push, a disc shell. In *In the Proceedings of the 2009 Principles of Distributed Computing Conference*, 2009.
- [17] E. Van Hensbergen and L. Ionkov. The v9fs project. <http://v9fs.sourceforge.net>.
- [18] E. Van Hensbergen and R. Minnich. Grave robbers from outer space using 9p2000 under linux. In *In Freenix Annual Conference*, pages 83–94, 2005.
- [19] L. Ionkov and E. Van Hensbergen. Xcpu2: Distributed seamless desktop extension. 2009.
- [20] Latchesar Ionkov, Ron Minnich, and Andrey Mirtchovski. The xcpu cluster management framework. In *First International Workshop on Plan9*, 2006.
- [21] R. Minnich and A. Mirtchovski. Xcpu: a new, 9p-based, process management system for clusters and grids. *Cluster Computing, 2006 IEEE International Conference on*, pages 1–10, 2006.
- [22] R. M. Needham and A. J. Herbert. *The Cambridge Distributed Computing System*. Addison-Wesley Publishers Limited, London, 1982.
- [23] R. Pike et al. Plan 9 from Bell Labs. *Computing Systems*, 8(3):221–254, Summer 1995.
- [24] Rob Pike, Dave Presotto, Ken Thompson, Howard Trickey, and Phil Winterbottom. The use of name spaces in plan 9. *SIGOPS Oper. Syst. Rev.*, 27(2):72–76, 1993.
- [25] Andrew S. Tanenbaum and Robbert Van Renesse. Distributed operating systems. *ACM Comput. Surv.*, 17(4):419–470, 1985.
- [26] IBM Blue Gene Team. Overview of the IBM BlueGene/P project. 2008.
- [27] Douglas Thain, Todd Tannenbaum, and Miron Livny. Distributed computing in practice: the condor experience. *Concurrency - Practice and Experience*, 17(2-4):323–356, 2005.
- [28] Vitanuova. Inferno operating system.
- [29] Y. Yu et al. DryadLINQ: A system for general-purpose distributed data-parallel computing using a high-level language. In *Symposium on Operating System Design and Implementation (OSDI), San Diego, CA, December*, pages 8–10, 2008.